

DOCUMENT: NovaCrest_CIM_2026.txt

TYPE: CIM

DATE: February 12, 2026

EXECUTIVE SUMMARY

NovaCrest Analytics (the "Company") is a premier B2B SaaS provider of financial data infrastructure, specifically engineered for banks, asset managers, and insurance companies. Founded four years ago by a team of former Goldman Sachs and Palantir executives, NovaCrest has rapidly emerged as the definitive system of record for real-time financial data normalization and portfolio risk analytics.

The Company is seeking an \$85M growth equity investment in exchange for a 30% minority stake, implying a pre-money valuation of \$283M. This capital will be deployed to accelerate European expansion (specifically a new London headquarters), build out a newly conceptualized predictive analytics product line, and scale the global headcount from 120 to 200 employees over the next 18 months.

NovaCrest represents a rare opportunity to back a highly profitable, founder-led Series B company operating in a massive, regulation-driven total addressable market. The Company has demonstrated exceptional financial performance, scaling revenue from \$8M in 2023, to \$18M in 2024, and ending 2025 with \$31M in Annual Recurring Revenue (ARR). The business operates with elite unit economics, boasting 78% gross margins, a 135% Net Retention Rate (NRR), and a highly efficient CAC payback period of 14 months. With a Rule of 40 score of 54 (72% YoY growth minus 18% Free Cash Flow margin due to aggressive reinvestment), NovaCrest is a category-defining asset.

COMPANY OVERVIEW

NovaCrest Analytics was founded in 2022 by Marcus Vance (CEO, former VP of Quantitative Trading at Goldman Sachs) and Elena Rostova (CTO, former Forward Deployed Engineer at Palantir). During his tenure at Goldman Sachs, Vance recognized a systemic inefficiency that was quietly draining profitability across Wall Street: financial institutions were spending billions of dollars and millions of highly skilled engineering hours simply cleaning, normalizing, and reconciling disparate data feeds before any actual trading or risk analysis could occur.

The mission of NovaCrest is to serve as the unified data translation layer for global finance. The Company's proprietary technology stack leverages a microservices architecture built on Kubernetes, utilizing a highly optimized proprietary data ingestion engine. This engine standardizes unstructured and semi-structured financial data—ranging from Bloomberg pricing feeds to obscure alternative data sets and legacy mainframe CSVs—into a single, queryable schema in milliseconds. By eliminating the "janitorial" data work, NovaCrest allows quantitative analysts and portfolio managers to focus exclusively on alpha generation and risk mitigation.

TECHNOLOGY PLATFORM

At the core of NovaCrest's competitive moat is its proprietary ontological translation engine, architected largely by Elena Rostova. Unlike legacy systems that rely on rigid, brittle ETL (Extract, Transform, Load) pipelines that break whenever a vendor changes an API endpoint, NovaCrest utilizes an adaptive schema layer.

This infrastructure handles the complexities unique to financial data: mapping different

identifiers (CUSIPs, ISINs, SEDOLs) to single corporate entities, reconciling asynchronous corporate actions (like stock splits or special dividends) across different time zones, and normalizing timestamp data to microsecond precision. The platform is entirely cloud-native, hosted primarily on AWS with multi-cloud redundancy, and is deployed directly into the client's Virtual Private Cloud (VPC) to ensure strict adherence to data sovereignty and privacy regulations.

MARKET OPPORTUNITY

The market for financial data infrastructure and Regulatory Technology (RegTech) is experiencing unprecedented macroeconomic tailwinds. The Total Addressable Market (TAM) for financial data management software is estimated at \$34 billion globally, with a Serviceable Addressable Market (SAM) of \$12 billion in North America and Western Europe.

Financial institutions are facing a dual mandate. First, they must modernize their legacy, on-premise tech stacks to remain competitive and to lay the groundwork for adopting Generative AI; AI models are useless if trained on fragmented, siloed data swamps. Second, they are simultaneously forced to comply with increasingly stringent global regulatory frameworks. The transition to T+1 settlement in the US, alongside the rollout of the Digital Operational Resilience Act (DORA) and MiFID II updates in Europe, is forcing firms to abandon manual data reconciliation.

The competitive landscape remains surprisingly fragmented. Legacy providers like Moody's Analytics and Regnology offer robust compliance tools but suffer from outdated, rigid architectures that require months of professional services to implement. Horizontal cloud data platforms like Snowflake Financial Services provide powerful raw storage and computing capability, but they are essentially empty warehouses—they require institutions to hire expensive internal engineering teams to build the actual financial data models. NovaCrest sits perfectly in the middle, offering cloud-native flexibility with out-of-the-box financial schema intelligence.

PRODUCT SUITE

NovaCrest operates a highly successful land-and-expand model centered around three core modules, with a fourth in active development:

1. DATABRIDGE

DataBridge is the Company's flagship product and the primary entry point for new customers. It is an API-based data normalization layer that connects to over 400 financial data providers (Bloomberg, Refinitiv, FactSet, IHS Markit, various alternative data vendors) and internal bank systems. It automatically cleanses, normalizes, and maps this data into a unified, strictly typed schema.

2. RISKLENS

RiskLens sits on top of DataBridge and provides real-time portfolio risk analytics. While legacy risk systems process overnight batches, RiskLens allows asset managers to run complex Monte Carlo simulations, multi-factor stress tests (e.g., simulating geopolitical shocks or interest rate hikes), and intraday liquidity analyses across multi-asset portfolios in real-time, directly leveraging the normalized data from DataBridge.

3. COMPLIANCE SYNC

ComplianceSync automates the most painful regulatory reporting workflows. It

automatically maps internal portfolio data to the specific reporting templates required by the SEC, FINRA, and European regulators. By generating audit-ready reports in minutes rather than weeks, it drastically reduces the manual overhead and risk of human error associated with quarterly and annual filings.

4. PREDICTIVE ANALYTICS (In Development)

A portion of the current fundraise is dedicated to launching a machine-learning-driven predictive analytics suite that will detect anomalies in trading patterns, forecast liquidity crunches based on historical cross-asset correlations, and identify alpha-generating signals from alternative data streams.

GO-TO-MARKET

NovaCrest employs a high-touch, direct enterprise sales motion targeting Chief Data Officers (CDOs), Chief Risk Officers (CROs), and VPs of Engineering at tier-one and tier-two financial institutions. The average sales cycle ranges from 6 to 9 months. This duration is primarily dictated by the stringent InfoSec and procurement requirements of the banking sector, requiring extensive Proof of Concept (PoC) testing and security audits.

The Company has secured highly lucrative, multi-year contracts with leading institutions, including Avalon Asset Management, Veracity Capital, Orion Global, and Zenith Partners. The "land and expand" strategy is highly effective: clients typically land with DataBridge for a specific trading desk or localized fund (Average ACV: \$180k). Within 12 to 18 months, they typically expand to RiskLens and ComplianceSync across the entire enterprise, driving ACVs upward of \$500k and sustaining the elite 135% Net Retention Rate.

FINANCIAL PERFORMANCE

NovaCrest has exhibited best-in-class financial metrics, transitioning rapidly from an early-stage startup to a highly efficient growth machine:

- 2023 ARR: \$8.0M
- 2024 ARR: \$18.0M
- 2025 ARR: \$31.0M
- Gross Margin: 78% (Optimized cloud architecture and containerization minimizes compute waste, a significant advantage over competitors who simply pass AWS costs onto the client).
- Net Retention Rate: 135% (Driven almost entirely by seat expansion within existing accounts and cross-selling the RiskLens module).
- CAC Payback: 14 months (Highly efficient for enterprise SaaS, driven by targeted outbound and strong industry word-of-mouth).
- LTV/CAC Ratio: 6.2x (Indicating substantial room to increase sales and marketing spend without deteriorating fundamental unit economics).
- Cost Structure: R&D accounts for 40% of revenue, Sales & Marketing at 25%, and G&A at 15%, leaving an 18% burn/investment margin.

MANAGEMENT TEAM

- Marcus Vance, CEO: Founder. Formerly VP of Quantitative Trading at Goldman Sachs, where he led a \$4B multi-strategy quantitative fund and directly managed the internal data engineering teams.
- Elena Rostova, CTO: Founder. Formerly Forward Deployed Engineer at Palantir, where she

designed the core data ontology frameworks for major European central banks and intelligence agencies.

- Robert Hughes, VP of Sales: Brought on in early 2024 to professionalize and scale the GTM motion. Formerly Head of Enterprise Sales at Bloomberg, overseeing a \$250M ARR book of business.
- Sarah Jenkins, CFO: Formerly Managing Director in the Tech/SaaS practice at Deloitte, specializing in scaling growth-stage software companies toward IPO readiness.

USE OF PROCEEDS

The \$85M growth equity injection will be utilized as follows:

- 40% (\$34M): European Expansion. Opening a London headquarters to capture the EMEA market. This requires establishing localized EU data centers, hiring specialized European compliance officers, and building a local GTM team to target institutions reacting to the Digital Operational Resilience Act (DORA) and MiFID II regulations.
- 35% (\$30M): Product Development. Launching the new AI-native predictive analytics product line, accelerating the machine learning roadmap, and adding an additional 100 native integrations to DataBridge.
- 25% (\$21M): Headcount Expansion. Scaling the organization from 120 to 200 employees over the next 18 months. This expansion is heavily weighted toward enterprise customer success managers and forward-deployed integration engineers to shorten client onboarding times from 4 months down to 45 days.

RISK FACTORS & MITIGANTS

- Customer Concentration: The Company's top three customers (Veracity Capital, Orion Global, and Zenith Partners) currently account for 40% of total ARR.
 - Mitigant: The new funding will expand the sales motion into mid-market asset managers and European banks, naturally diluting this concentration. No single customer represents more than 15%.
- Regulatory Dependency: A portion of the value proposition for ComplianceSync relies on the complexity of current regulations. Simplification of these rules could reduce product stickiness.
 - Mitigant: Regulatory complexity in finance historically only moves in one direction: increasing complexity. Furthermore, ComplianceSync is only one of three core modules.
- Implementation Bottlenecks: Historical deployments have stretched to 3-4 months, occasionally leading to churn (e.g., Cornerstone Bank) due to scope creep and lack of customer success resources.
 - Mitigant: The use of proceeds specifically addresses this by allocating capital to build out a robust, proactive customer success and implementation engineering team.
- Key Person Risk: The proprietary data ingestion architecture was largely architected by CTO Elena Rostova.
 - Mitigant: The engineering leadership team has been significantly expanded, the codebase has been aggressively documented, and core architecture responsibilities have been successfully distributed among a team of five senior

engineering directors.

===== END OF DOCUMENT =====

DOCUMENT: NovaCrest_Mgmt_Presentation_Transcript.txt

TYPE: Management Presentation

DATE: February 25, 2026

SPEAKER: JASON REED (Partner, Meridian Growth Partners)

SPEAKER: CHLOE ADAMS (Associate, Meridian Growth Partners)

SPEAKER: MARCUS VANCE (CEO, NovaCrest Analytics)

SPEAKER: ELENA ROSTOVA (CTO, NovaCrest Analytics)

JASON REED: Good morning, Marcus, Elena. Thanks for having us. The Meridian team has been digging into the CIM, the preliminary data room, and the financial model you sent over on Tuesday. Let me just start by saying that the growth you've charted over the last three years is incredibly impressive. Today, we really want to unpack the story behind the numbers, get under the hood of the technology, and frankly, talk about the path to \$100M ARR.

MARCUS VANCE: Absolutely, Jason. Great to see you and Chloe again. I know we've got a lot to cover, so let's dive right in. The origin story of NovaCrest really starts on the trading floor at Goldman. I was managing a massive quant desk, and I realized that my most expensive, PhD-level quantitative analysts—guys getting paid seven figures to generate alpha—were spending eighty percent of their day playing data janitor. They were writing bespoke ETL pipelines to stitch together Bloomberg pricing data, alternative data streams like satellite imagery, and our own legacy internal risk metrics. It was a nightmare. Every single bank does this. They all build these fragile, duct-taped in-house systems that completely break the second a vendor changes a single API endpoint or alters a CUSIP format.

I partnered with Elena, who was solving similar, massive-scale data ontology problems at Palantir for intelligence agencies, and we built DataBridge. We aren't just selling a SaaS tool; we are selling the fundamental infrastructure that allows banks to actually use their data. We are laying the digital railroad tracks.

ELENA ROSTOVA: Exactly. To expand on Marcus's point, from a technical architecture standpoint, what we've built is a deterministic translation layer. When you look at legacy systems in finance, they usually force the bank to conform to the vendor's rigid data model. DataBridge flips that dynamic entirely. We built an adaptive ontology framework. We ingest the messy, unstructured, chaotic reality of the financial markets and map it to a unified, mathematically rigorous schema inside the client's environment. Our true data moat isn't just the software interface; it's the 400-plus custom vendor integrations we've actively maintained and updated over the last four years. If Refinitiv changes a field name on a bond pricing feed at 3:00 AM, our system adapts dynamically, and the portfolio manager's risk dashboard never goes down.

CHLOE ADAMS: Elena, on that note of stickiness, we see the 135% Net Retention Rate, which is fantastic and puts you in the top quartile of SaaS metrics. But digging into the customer tape, we did notice you lost two mid-sized clients last year, one of them being Cornerstone Bank. I'd love to hear your perspective on what happened there. Is that an isolated incident, or a symptom of a broader onboarding bottleneck?

MARCUS VANCE: I can take that one, Chloe. Look, I'll be completely transparent. When you are

scaling an enterprise company from \$8 million to \$31 million in two years, you learn exactly who your ideal customer profile is, and just as importantly, who it isn't. Cornerstone came on right after our Series A. They are a great mid-market bank, but they completely lacked the internal engineering maturity to utilize our platform. They basically wanted us to act as an outsourced IT consulting firm and build their entire cloud data warehouse for them from scratch. That's not our model. We are a software company, not Accenture. We had massive scope creep during the 4-month integration phase, our support team was stretched thin dealing with their legacy on-premise servers, and at renewal, we made a mutual decision to part ways rather than compromise our gross margins doing custom professional services. It was a painful, but necessary, learning experience.

JASON REED: Got it. I appreciate the candor. Let's talk about the competitive landscape. Bloomberg recently announced a massive push into enterprise data management with their B-PIPE expansion. How do you defend against a juggernaut that already sits on virtually every trader's desk in the world?

MARCUS VANCE: It's a great question, Jason, and one we get often. Bloomberg's strategy is fundamentally about locking you into the Bloomberg ecosystem. They want to be a walled garden. If you use their data management platform, it works beautifully—as long as you only use Bloomberg data. But the minute you try to pipe in proprietary alternative data, or competitor pricing feeds, or legacy internal risk metrics, the friction is massive by design.

NovaCrest is Switzerland. We are entirely agnostic. We sit above the data providers. Furthermore, our RiskLens module processes computations in real-time, in milliseconds. The legacy providers are still largely relying on overnight batch processing, which simply cannot match our latency. In today's volatile markets, overnight risk reporting is useless.

CHLOE ADAMS: Elena, regarding the engineering team. You're scaling fast. We flagged some potential key person risk in the CIM. How is engineering retention holding up, and how much of this architecture is reliant specifically on you being in the weeds every day?

ELENA ROSTOVA: It's a fair point, Chloe. In the first two years, I personally wrote a significant portion of the core ingestion engine. But over the last 18 months, my role has shifted entirely to architecture and team building. We've hired phenomenal talent, primarily pulling senior engineers from Palantir, Stripe, and Two Sigma. Engineering retention is currently at 94%, well above industry average. The architecture is now fully modularized, aggressively documented, and distributed across five senior engineering directors. The primary reason we are raising this \$85M is to aggressively expand this team. We are going from 120 to 200 people, and a significant portion of that will be forward-deployed integration engineers. That directly solves the Cornerstone issue—we will have the dedicated headcount to hold the client's hand through the complex legacy integrations without burning out our core product engineers.

JASON REED: Let's transition to the future. To get to \$100M ARR, you need to execute on this European expansion and launch this new predictive analytics suite. What makes you confident you can win in London against entrenched local players?

MARCUS VANCE: It comes down to one word: Regulation. The digital compliance landscape in Europe is vastly more complex and punitive than in the US right now. The Digital Operational Resilience Act, or DORA, is forcing every European asset manager to completely audit, stress-test, and localize their third-party data providers. By opening our London office and

establishing localized EU data centers, we can offer ComplianceSync as an out-of-the-box solution for DORA and MiFID II. European banks are terrified of vendor lock-in and terrified of non-compliance fines. We already have three letters of intent from tier-one European asset managers, strictly contingent on us standing up localized support and SOC 2 Type II certification. That expansion alone is a \$40 million revenue opportunity over the next three years. That's the path to \$100 million.

===== END OF DOCUMENT =====

DOCUMENT: Expert_Call_DataInfra_Specialist.txt

TYPE: Expert Call

DATE: March 2, 2026

SPEAKER: CHLOE ADAMS (Associate, Meridian Growth Partners)

SPEAKER: DR. SARAH CHEN (Independent Consultant, Former Chief Data Officer at EuroBank Corp)

CHLOE ADAMS: Dr. Chen, thank you so much for taking the time today. To set the stage, we are looking at a potential growth equity investment in a company called NovaCrest Analytics. As you know, they provide an API-based data normalization and risk analytics platform for financial institutions. Given your extensive background as a Chief Data Officer, we'd love your unvarnished take on this market. Do banks actually need third-party tools for this, or are they just building it themselves internally?

DR. SARAH CHEN: Hi Chloe, happy to be here. The short answer is: yes, they absolutely need it. The days of banks building their own data pipelines from scratch are ending. It is a fool's errand. When I was the CDO at EuroBank, we had a team of fifty engineers—very expensive engineers—just writing ETL scripts to parse data from Refinitiv, Bloomberg, and our own legacy mainframe systems. It was a nightmare.

To give you an example, just mapping CUSIPs to ISINs accurately across different international asset classes, while accounting for corporate actions like stock splits, is brutally complex. Every time a vendor updated an API, or changed a comma in a CSV export, our internal systems broke and the risk dashboard for the C-suite went dark. The TAM for what NovaCrest is doing is enormous because every bank on the planet is trying to figure out how to structure their data for AI initiatives. You simply can't do Generative AI or predictive modeling if your underlying data is a disorganized swamp.

CHLOE ADAMS: That makes a lot of sense. How do you view NovaCrest compared to the alternatives out there? For example, we see a lot of institutions moving their entire data lakes to Snowflake Financial Services, or using legacy regulatory players like Regnology and Moody's Analytics. Where does NovaCrest fit?

DR. SARAH CHEN: They serve very different purposes. Snowflake is phenomenal. It's an incredible database architecture. But Snowflake is basically an empty, highly efficient warehouse. You still have to build the shelves, organize the boxes, write the labels, and figure out the logistics. Snowflake Financial Services gives you some industry-specific templates, but you still need a heavy lift from data engineering to make it work for complex derivatives or multi-asset portfolio risk. You are paying for compute and storage, not financial intelligence. Regnology and Moody's are the opposite. They are very structured, very good for regulatory compliance, but they are rigid and largely built on legacy architectures. If you want to run a

custom Monte Carlo simulation on an obscure, illiquid asset class, you are going to have a bad time with a legacy provider.

NovaCrest is highly regarded right now because they provide the flexibility and speed of a modern cloud stack with the out-of-the-box financial intelligence of a legacy provider. Their DataBridge product is excellent at translating the ontology of finance seamlessly.

CHLOE ADAMS: Have you worked directly with NovaCrest, or evaluated them in a formal procurement process?

DR. SARAH CHEN: Yes, I evaluated them about fourteen months ago during a consulting engagement for a mid-sized asset manager in London. We ultimately selected them over Charles River, but I will give you a candid negative. Their onboarding process was incredibly painful. It took us 4 months to get DataBridge fully integrated into our legacy on-premise systems. That's too long for most banks. It wasn't that their software was broken; it was that mapping our highly customized internal bond fields to their schema took endless back-and-forth. Once it was up and running, it was flawless, but their deployment team was clearly stretched too thin. They simply didn't have the headcount or the technical account managers to hold the client's hand through the process.

CHLOE ADAMS: That aligns with some of our other findings, and it's something the management team has acknowledged. They are actually planning to use a large portion of this funding to expand into Europe and build out that deployment team. As someone based in London, how viable is their European expansion strategy?

DR. SARAH CHEN: Extremely viable, but they have to navigate the regulatory environment carefully. Under DORA and the incoming MiFID II updates, operational resilience and data sovereignty are paramount. European banks are terrified of vendor lock-in, and they are legally prohibited from using non-EU data hosting for certain client data. If NovaCrest is bringing \$30 or \$40 million to establish local data centers in Frankfurt and London, and to hire local compliance experts who speak the language of European regulators, they will win deals. The demand is absolutely there, but they cannot service Europe successfully from a timezone in New York.

===== END OF DOCUMENT =====

DOCUMENT: Expert_Call_Former_Customer_VP_Eng.txt

TYPE: Expert Call

DATE: March 5, 2026

SPEAKER: JASON REED (Partner, Meridian Growth Partners)

SPEAKER: JAMES PARK (Former VP of Engineering, Cornerstone Bank)

JASON REED: James, thanks for jumping on the phone today. We're currently looking closely at a potential investment in NovaCrest Analytics. We noticed in their data room that Cornerstone Bank was a client, but you churned last year. When we brought this up, the CEO of NovaCrest suggested that Cornerstone simply wasn't ready for an enterprise-scale integration, and that the deal suffered from massive scope creep. We want to hear your unvarnished side of the story.

JAMES PARK: [Laughs] Well, I suppose that's one way for Marcus to spin it. I appreciate the candor, Jason. Look, I will start by saying something that might surprise you: the core product is genuinely best-in-class. Elena and her engineering team have built a miraculous piece of

technology. When we first did the Proof of Concept with DataBridge in 2023, we tested it against our messiest fixed-income and structured credit data. It mapped the data, resolved the identifier conflicts, and presented a clean schema faster than anything we had ever seen. The tech works.

JASON REED: So if the tech works, what went wrong? Why did you rip it out?

JAMES PARK: Two things happened simultaneously: a complete support collapse and pricing arrogance. We signed with them right around the time they raised their Series A. Before that raise, if we had a bug or an integration hurdle, Elena or one of her lead engineers would be on a Slack huddle with us in ten minutes. After the Series A, they aggressively scaled their sales team to chase new logos, but they clearly didn't scale their customer success or technical support teams to match.

We did run into some complex scope creep during implementation—which, to be fair to Marcus, was partially our fault because our legacy SQL database was a disorganized mess. But their support response times degraded severely. A critical API routing ticket for our end-of-day reconciliation took 4 days to resolve at one point. In institutional banking, a 4-day data outage on a risk dashboard during quarter-end is a fireable offense. We were flying blind.

JASON REED: That's a serious operational risk. What about the pricing issue you mentioned?

JAMES PARK: We were originally on a \$250k annual contract. When the renewal came up in late 2024, Robert Hughes, their new VP of Sales, hit us with a 40% price jump. He claimed it was to cover the "custom integration maintenance and excessive compute" we required due to our legacy systems. We balked. My CFO nearly threw him out of the room. It felt like a bait-and-switch. We were already frustrated with the support, so we decided not to renew.

JASON REED: What did you switch to, and how is it performing compared to NovaCrest?

JAMES PARK: We migrated our entire data lake to Snowflake and decided to build the translation layer in-house using an open-source Python framework. Between you and me? It cost us a fortune in engineering hours, my engineers hated doing the janitorial data mapping, and honestly, it doesn't run nearly as smoothly or as fast as NovaCrest did. We traded a vendor problem for a massive internal maintenance burden.

JASON REED: Given that, what would NovaCrest need to fix to win a bank like Cornerstone back?

JAMES PARK: Their go-to-market and post-sales support just weren't ready for enterprise scale. They treated a complex banking integration like a self-serve SaaS product. If they come back to me tomorrow with a dedicated technical account manager, guaranteed Service Level Agreements (SLAs) with financial penalties for response times, and a transparent pricing model that doesn't penalize us for integration hurdles, I would actually take the meeting. As I said, the tech is that good. The business operations around the tech were just immature. If they use your money to fix the operations, they have a monster business.

===== END OF DOCUMENT =====

DOCUMENT: Customer_Reference_Avalon_AM.txt

TYPE: Customer Reference

DATE: March 8, 2026

SPEAKER: CHLOE ADAMS (Associate, Meridian Growth Partners)

SPEAKER: LISA TORRES (Head of Data Engineering, Avalon Asset Management)

CHLOE ADAMS: Hi Lisa, thank you so much for making the time today. We are conducting due diligence on NovaCrest Analytics. As a current, high-value customer managing over \$1.2B in AUM, your perspective is incredibly valuable to our deal team. Why did Avalon choose NovaCrest over building the infrastructure internally or using legacy platforms?

LISA TORRES: Hi Chloe. Sure thing, I'm happy to help. We looked at absolutely everything on the market before deciding. We evaluated Palantir Foundry, we looked at building a custom solution on top of AWS, and we looked at legacy portfolio management systems like Charles River.

NovaCrest won primarily because of time-to-value. Palantir Foundry is amazing—it's like a Ferrari—but it requires an army of consultants to run, and it was prohibitively expensive for our AUM tier. Building internally was out of the question; we are a lean team at Avalon. We don't have fifty data engineers to throw at a problem. DataBridge was the only solution that connected to our FactSet terminals, our internal prime brokerage feeds, and our alternative data sources out of the box. We were normalized and querying data within weeks, not years. It just works.

CHLOE ADAMS: That's a great endorsement. Can you quantify the Return on Investment you've seen since the implementation?

LISA TORRES: Absolutely. Before NovaCrest, my team spent the first four hours of every single trading day just reconciling data discrepancies. We had portfolio managers screaming because their P&L didn't match the risk system due to asynchronous pricing feeds. It was miserable, high-stress work. NovaCrest automated that entire pipeline. We reduced data reconciliation time by 70%. From a hard cost perspective, I estimate we've saved about \$2 million a year in manual processing costs, data errors, and avoided having to hire at least five additional senior data engineers.

CHLOE ADAMS: How has the reliability and support been? We've heard some mixed reviews from early-stage clients regarding onboarding support and ticket response times.

LISA TORRES: We haven't experienced those issues, but keep in mind we signed with them in late 2024. By that point, I think they had matured their processes significantly compared to their early days. The platform reliability is rock solid. We track them closely, and they have maintained five nines of uptime. Their enterprise support desk has been highly responsive to us, usually resolving mapping questions within the hour.

CHLOE ADAMS: Are you looking to expand your usage of the platform, or are you just maintaining the current DataBridge deployment?

LISA TORRES: We are absolutely expanding. We currently use DataBridge globally across all our desks, and we are running an active pilot program right now to add the RiskLens module next quarter. RiskLens is going to allow our quantitative team to run multi-factor stress tests in real-time. For example, we can simulate a geopolitical shock in Taiwan and see the immediate liquidity impact across our entire equities and derivatives portfolio in about 400 milliseconds. That is a game changer for our macro strategy and justifies the upsell cost entirely.

CHLOE ADAMS: Are there any concerns or red flags you would point out to an investor? Anything they are missing?

LISA TORRES: There is one major concern, and it's an administrative and compliance one. We are dealing with massive institutional capital, and our compliance and vendor risk management

teams require rigorous audits of our critical software vendors. Right now, NovaCrest only has a SOC 2 Type I certification, which basically just says they have policies written down. We need to see SOC 2 Type II, which proves those policies are actively followed over a prolonged period. I've brought this up with Marcus and Robert Hughes repeatedly. They promised it is on the roadmap for this year, but our compliance team is breathing down my neck about it and threatening to halt the RiskLens expansion. If they want to close bigger enterprise deals, or keep highly regulated institutions like us happy long-term, they absolutely have to secure that SOC 2 Type II certification immediately.

CHLOE ADAMS: That is incredibly helpful insight, Lisa. We will definitely bring that up with the management team. Thank you so much for your time today.

===== END OF DOCUMENT =====